

Ex1.StdDev

January 16, 2024

1 1st exercise: Work with standard deviations for anomaly detection

- Course: AML
- Lecturer: Gernot Heisenberg
- Author of notebook: Finn Heydemann
- Date: 28.10.2023

GENERAL NOTE 1: Please make sure you are reading the entire notebook, since it contains a lot of information on your tasks (e.g. regarding the set of certain parameters or a specific computational trick), and the written mark downs as well as comments contain a lot of information on how things work together as a whole.

GENERAL NOTE 2: * Please, when commenting source code, just use English language only. * When describing an observation please use English language, too * This applies to all exercises throughout this course.

1.0.1 DESCRIPTION:

This notebook allows you for getting into standard deviations as a common technique to detect anomalies when the data is normally distributed.

1.0.2 TASKS:

The tasks that you need to work on within this notebook are always indicated below as bullet points. If a task is more challenging and consists of several steps, this is indicated as well. Make sure you have worked down the task list and commented your doings. This should be done by using markdown. Make sure you don't forget to specify your name and your matriculation number in the notebook.

YOUR TASKS in this exercise are as follows: 1. import the notebook to Google Colab or use your local machine. 2. make sure you specified your name and your matriculation number in the header below my name and date. * set the date too and remove mine. 3. read the entire notebook carefully * add comments wherever you feel it necessary for better understanding * run the notebook for the first time. * understand the output 4. go and find three different data sets on the web * kaggle.com might be a good source (they also offer an API for data download) * make sure two of the three data sets are normally distributed * download one data set that is not normally distributed 5. visualize the data 6. compute the anomalies 7. visualize the anomalies 8. does the 0,3% rule apply? 9. what are differences between the normally distributed and the non-normally distributed data sets with respect to the outlier detection? 10. which statement can be made and which cannot? _____

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from numpy.random import randn
np.random.seed(1)
```

```
[2]: # Function to Detection Outlier on one-dimentional datasets.
def find_anomalies(random_data):
    #define a list to accumlate anomalies
    anomalies = []

    # Set upper and lower limit to 3 standard deviation
    random_data_std = np.std(random_data)
    random_data_mean = np.mean(random_data)
    anomaly_cut_off = random_data_std * 3

    lower_limit = random_data_mean - anomaly_cut_off
    upper_limit = random_data_mean + anomaly_cut_off

    print("lower limit=", round(lower_limit,8))
    print("upper limit=", round(upper_limit,8))

    # Generate outliers list
    for outlier in random_data:
        if outlier > upper_limit or outlier < lower_limit:
            anomalies.append(outlier)

    return anomalies
```

```
[3]: # multiply and add by random numbers to get some real values
# randn generates samples from the normal distribution (important - see below)
data = np.random.randn(50000) * 20 + 20
```

```
[4]: anomalies = find_anomalies(data)
```

lower limit= -39.96482266

```
upper limit= 80.11636225
```

1.1 Result

These anomalies are exceeding the lower and upper 3rd scatter range. Thus, statistically spoken, they do belong to a population size of less than 0,3% of the entire data set! For sure, the above conclusion is true if and only if the data is normally distributed!

```
[5]: print(anomalies)
```

```
[80.61714224744061, 99.17205408075927, 86.4215751234073, -41.0752876085261,
-43.06714900382097, -40.32063970418597, 88.65326863591119, -41.28282712210201,
94.80497807409178, 82.70094680175814, 84.76686393504752, -45.06068469953237,
82.64059456281011, 80.25509136474466, -43.81232693648161, -45.21230191553079,
88.08604551497943, 82.369591815651, 85.9708104656157, 82.97968050439259,
-45.8971681339703, -41.11266480828613, 100.53698089094755, -44.38264211247683,
82.23783662033449, 92.26554014166295, 87.96231313004853, 91.21746641411097,
-53.128801985095905, -45.641575926212354, 91.21225295344836, -43.46923274683489,
82.13498657145581, 82.90737927394908, -46.216851244091984, -41.9565310384887,
-45.6065519410858, 80.5243698860871, -43.30420249353775, -42.25263324372853,
81.59215546151304, 103.36235355910189, 96.68762041821407, -40.586879558462776,
-48.71851620008829, -45.18412137279293, 81.54158108441533, 84.5587224070789,
-40.7688517191754, -44.08128900692253, -42.043297873881315, 87.71430311862296,
88.9391203157937, -40.52113314693459, -49.02805812510347, 86.72050502651298,
-44.452278374625166, -58.55028132123577, 87.83885146048503, -41.810065776968784,
-42.38237155225082, 91.19937116220078, 93.19531683513023, -41.0172829464469,
82.94710559091223, 86.8410257826826, -47.358945718270846, -47.27471768342224,
-44.05283959700523, 81.39224116690686, -64.66329594864382, -41.91924535511986,
86.98293402121955, -42.78324538163221, -43.44070544512959, 83.14254869231635,
80.29722046012112, 80.97664481662183, -49.135306899574985, 90.95359817188209,
85.7504914101956, 90.63353726783707, -52.25624570621119, -48.427273980448774,
92.45148786479089, -40.60136533430394, 90.66197938200516, -40.652000332079666,
85.47272253055641, -40.59878231029168, -53.957837173283735, -40.80036485428288,
82.4740392321337, 82.69560613707037, -46.18148126636443, 86.9908487614501,
80.1457297060347, 82.24721888631221, -45.832275955612076, 84.176266941649,
84.4193809675114, -48.12080165824392, -47.66016305888587, 83.03299706691035,
-46.34572895708007, -43.94127801481702, -44.888893484963376, 82.73490952329395,
83.57335544403156, 80.7727413101268, -52.50413207267792, -49.69745467233162,
80.22051938742219, -61.76977168084022, -41.57467213920625, -58.86284640187061,
-51.197383857166216, -56.11928395638613, -48.85042290718998, -52.64159591899764,
-40.858015925037485, 90.00512514161373, -44.150803760611865, 81.0757508416524,
88.42082610751223, -48.97515466922074, -41.87410601743178, 81.91272478059072,
84.69303765209298, -40.93984641793596, -52.615615036883995, -49.39834453038375,
84.48319182306491, 80.41393381216217, 89.04319371214211, 83.11545348720793,
86.3948526034166, -54.34260433053814, -53.884089553553295, -60.83360878272147,
-45.502951473453436, -42.84117089742673]
```

```
[6]: data = pd.read_csv("data/exercise_1/StudentsPerformance.csv")
```

```
[7]: data.head()
```

```
[7]:  gender race/ethnicity parental level of education      lunch \
0  female      group B      bachelor's degree      standard
1  female      group C      some college      standard
2  female      group B      master's degree      standard
3   male      group A      associate's degree  free/reduced
4   male      group C      some college      standard


      test preparation course  math score  reading score  writing score
0                none          72          72          74
1          completed          69          90          88
2                none          90          95          93
3                none          47          57          44
4                none          76          78          75
```

```
[8]: import matplotlib.pyplot as plt

fig, axs = plt.subplots(3, figsize=(8, 12), layout='constrained', sharex=True)
fig.suptitle("Male and Female", fontsize=20)
for ax, key in zip(axs, ["math score", "writing score", "reading score"]):
    ax.hist(data[key], bins=20)
    ax.set_title(key)

plt.show()
```

Male and Female




```
[9]: fig, axs = plt.subplots(3, figsize=(8, 12), layout='constrained', sharex=True)
fig.suptitle("Male Only", fontsize=20)
for ax, key in zip(axs, ["math score", "writing score", "reading score"]):
    ax.hist(data[data["gender"] == "male"][key], bins=20)
    ax.set_title(key)


plt.show()
```

Male Only


math score



writing score



reading score




```
[10]: fig, axs = plt.subplots(3, figsize=(8, 12), layout='constrained', sharex=True)
fig.suptitle("Female Only", fontsize=20)
for ax, key in zip(axs, ["math score", "writing score", "reading score"]):
    ax.hist(data[data["gender"] == "female"][key], bins=20)
    ax.set_title(key)


plt.show()
```


Female Only


math score



writing score



reading score



```
[11]: data["math score"].mean(), data["math score"].std()
```

```
[11]: (66.089, 15.16308009600945)
```

```
[12]: def anom_detection(data: np.ndarray or pd.Series):  
        return data[np.logical_or(data < data.mean() - 3 * data.std(), data > data.  
        ↪mean() + 3 * data.std())]  
  
anom_detection(data["math score"])
```

```
[12]: 17      18  
59      0  
787     19  
980      8  
Name: math score, dtype: int64
```

```
[13]: ind = anom_detection(data["math score"]).index  
data.loc[ind]
```

```
[13]:      gender race/ethnicity parental level of education      lunch \  
17   female      group B      some high school  free/reduced  
59   female      group C      some high school  free/reduced  
787  female      group B      some college      standard  
980  female      group B      high school  free/reduced  
  
      test preparation course  math score  reading score  writing score  
17      none                18          32          28  
59      none                0          17          10  
787     none                19          38          32  
980     none                8          24          23
```

```
[14]: data = pd.read_csv("data/exercise_1/city_temperature.csv")
```

```
/tmp/ipykernel_30670/3924103419.py:1: DtypeWarning: Columns (2) have mixed  
types. Specify dtype option on import or set low_memory=False.  
data = pd.read_csv("data/exercise_1/city_temperature.csv")
```

```
[15]: data.head()
```

```
[15]:   Region Country State   City Month Day Year AvgTemperature  
0  Africa  Algeria  NaN  Algiers    1   1  1995          64.2  
1  Africa  Algeria  NaN  Algiers    1   2  1995          49.4  
2  Africa  Algeria  NaN  Algiers    1   3  1995          48.8  
3  Africa  Algeria  NaN  Algiers    1   4  1995          46.4  
4  Africa  Algeria  NaN  Algiers    1   5  1995          47.9
```

```
[16]: data.drop(data[data["AvgTemperature"] == -99].index, inplace=True)
```

```
[17]: data["AvgTemperature"] = (data["AvgTemperature"] - 32) * 5/9
```

```
[18]: data["date"] = pd.to_datetime(data[["Year", "Month", "Day"]])
data = data.set_index("date").groupby(["City", pd.
↳Grouper(freq="D")])["AvgTemperature"].mean()
data = data.reset_index()
```

```
[19]: data
```


```
[19]:
```

	City	date	AvgTemperature
0	Abidjan	1995-01-01	28.111111
1	Abidjan	1995-01-02	27.944444
2	Abidjan	1995-01-03	27.222222
3	Abidjan	1995-01-04	28.500000
4	Abidjan	1995-01-05	28.555556
...
2769402	Zurich	2020-05-09	19.500000
2769403	Zurich	2020-05-10	18.166667
2769404	Zurich	2020-05-11	11.111111
2769405	Zurich	2020-05-12	6.388889
2769406	Zurich	2020-05-13	7.000000


[2769407 rows x 3 columns]

```
[20]: citys = np.unique(data["City"])
fig, axs = plt.subplots(len(citys) // 10, 10, sharey=True, figsize=(15, 40),
↳tight_layout=True)
for ax, city in zip(axs.flatten(), citys):
    p_data = data[data["City"] == city]["AvgTemperature"]
    weights = np.ones_like(p_data)/float(len(p_data))
    ax.hist(p_data, bins=20, weights=weights)
    ax.set_title(city)

plt.show()
```



```
[21]: plt.hist(data[data["City"] == "Brasilia"]["AvgTemperature"], bins=20)
plt.show()
```



```
[22]: ind = anom_detection(data[data["City"] == "Brasilia"]["AvgTemperature"]).index
```


```
[23]: data.loc[ind]
```

```
[23]:
```

	City	date	AvgTemperature
406288	Brasilia	1996-06-29	13.666667
406289	Brasilia	1996-06-30	14.333333
406633	Brasilia	1997-06-09	15.111111
406771	Brasilia	1997-10-25	28.611111
406986	Brasilia	1998-06-02	15.500000
407080	Brasilia	1998-09-06	30.000000
407223	Brasilia	1999-02-17	28.500000
408563	Brasilia	2002-10-26	30.944444
409198	Brasilia	2004-07-24	15.277778
409548	Brasilia	2005-07-09	15.388889
409549	Brasilia	2005-07-10	14.333333
409550	Brasilia	2005-07-11	14.611111

409908	Brasilia	2006-07-04	15.388889
409919	Brasilia	2006-07-15	14.888889
410642	Brasilia	2008-07-09	14.888889
410643	Brasilia	2008-07-10	15.444444
410647	Brasilia	2008-07-14	15.333333
412910	Brasilia	2014-09-29	28.388889
412929	Brasilia	2014-10-18	28.777778
413914	Brasilia	2017-07-04	15.111111
413915	Brasilia	2017-07-05	15.055556
413916	Brasilia	2017-07-06	15.555556
414017	Brasilia	2017-10-15	29.000000
414262	Brasilia	2018-06-18	13.388889
414712	Brasilia	2019-09-20	28.222222
414713	Brasilia	2019-09-21	29.555556
414742	Brasilia	2019-10-20	28.555556
414765	Brasilia	2019-11-12	28.500000

```
[24]: plt.hist(data[data["City"] == "Pyongyang"]["AvgTemperature"], bins=20)
plt.show()
```



```
[25]: anom_detection(data[data["City"] == "Pyongyang"]["AvgTemperature"])
```

```
[25]: Series([], Name: AvgTemperature, dtype: float64)
```

Obviously the anomaly detection doesn't work when the data is not normally distributed. When checking the daily temperatures in Pyongyang no anomalies are detected. Whereas when the daily temperatures are evenly distributed, anomalies are detected. Same goes for the student performances.

```
[ ]:
```